

raelize

Acquisition in the billions: Breaking cryptographic keys with fast SCA

Cristofaro Mune

cristofaro@raelize.com

[@pulsoid](https://twitter.com/pulsoid)

Introduction.

Me

Cristofaro Mune

- Co-Founder at Raelize; Security Researcher
- 20+ years in security
- 15+ years analyzing the security of complex systems and devices
- Based in The Netherlands. Specialized in **Device Security**
- Security testing, Consultancy and Training
- Low level software, hardware security:
 - Secure Boot, TEE, Fault injection,...



Our research: <https://raelize.com/blog>

Goals

- Outline usage of cryptographic **keys** in **modern devices**
- Introduce **side-channel-attacks** (SCA)
- Breaking AES via **power analysis**
 - on a modern *System-on-Chip* (SoC)
- Demonstrate techniques for **fast** acquisition
 - **Billions** of traces per day
- Reflect on implications

Raise **awareness**

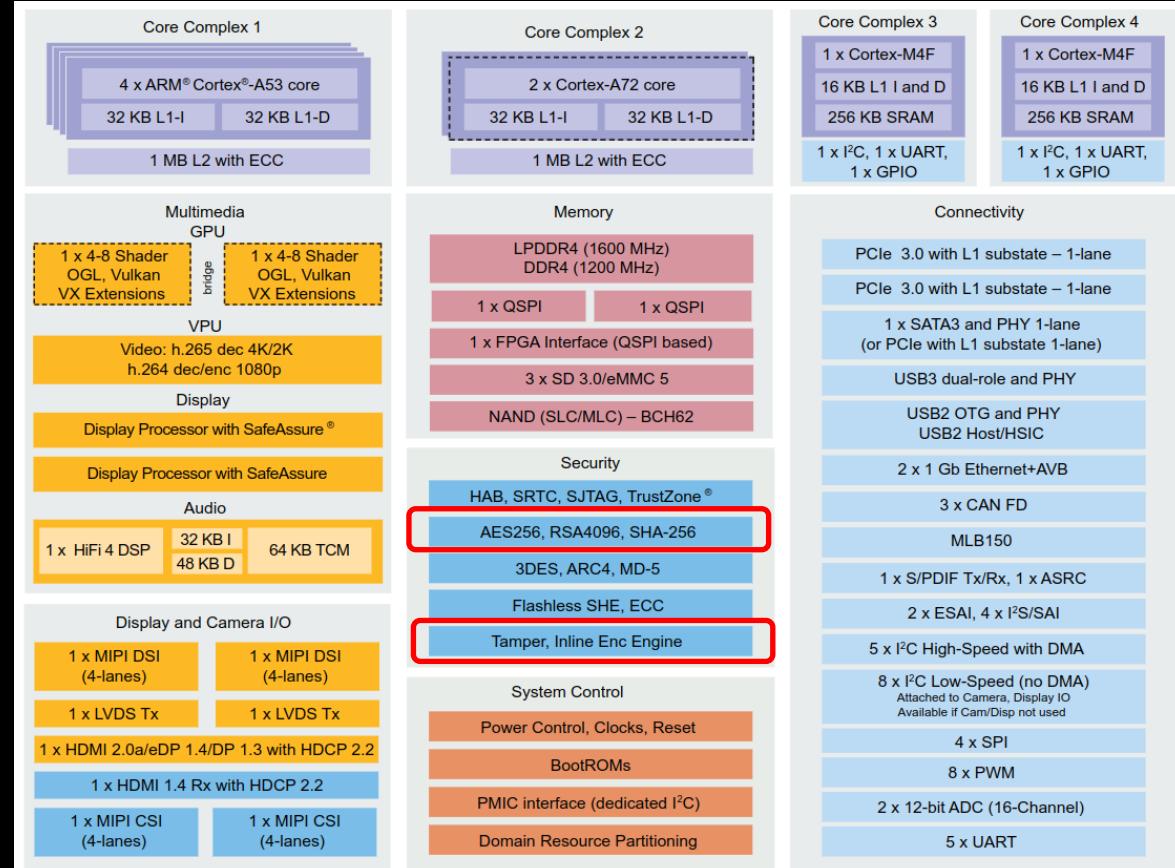
Of devices, keys and crypto(-graphy).

Devices: cryptographic operations

- Several designs (and implementations) available
- For our purposes, let's consider the following:
 - Pure **software** (SW):
 - Also in white-box cryptography (WBC) form
 - **Hardware-assisted:**
 - i.e. make use of hardware (HW) cryptographic **accelerators**

No intention of **completeness!**

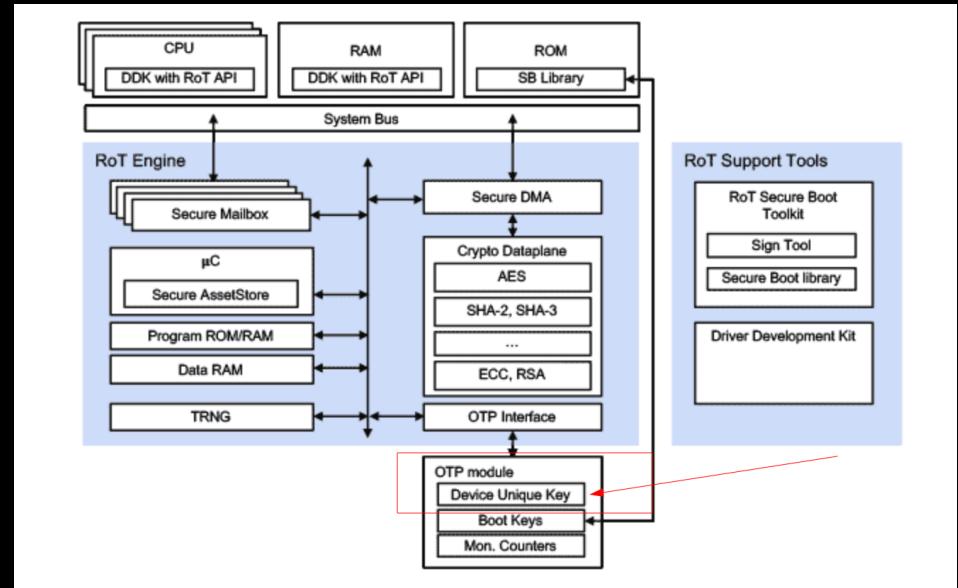
Modern device: HW cryptographic accelerators



NXP i.MX8 SoC

Hardware: Keys inaccessible to SW

- Device Unique Key(s):
 - Stored in (e-)fuses or in the actual digital logic (rarely)
- Directly loaded in HW crypto engines slots
- No way for SW to read such keys:
 - No interface available



Rambus - RT-260 Root of Trust

storage systems, which protects data from a compromised “global” key. Device binding is based on a per-device unique identity which is baked into the hardware and cannot be altered by software after the devices are initialized.

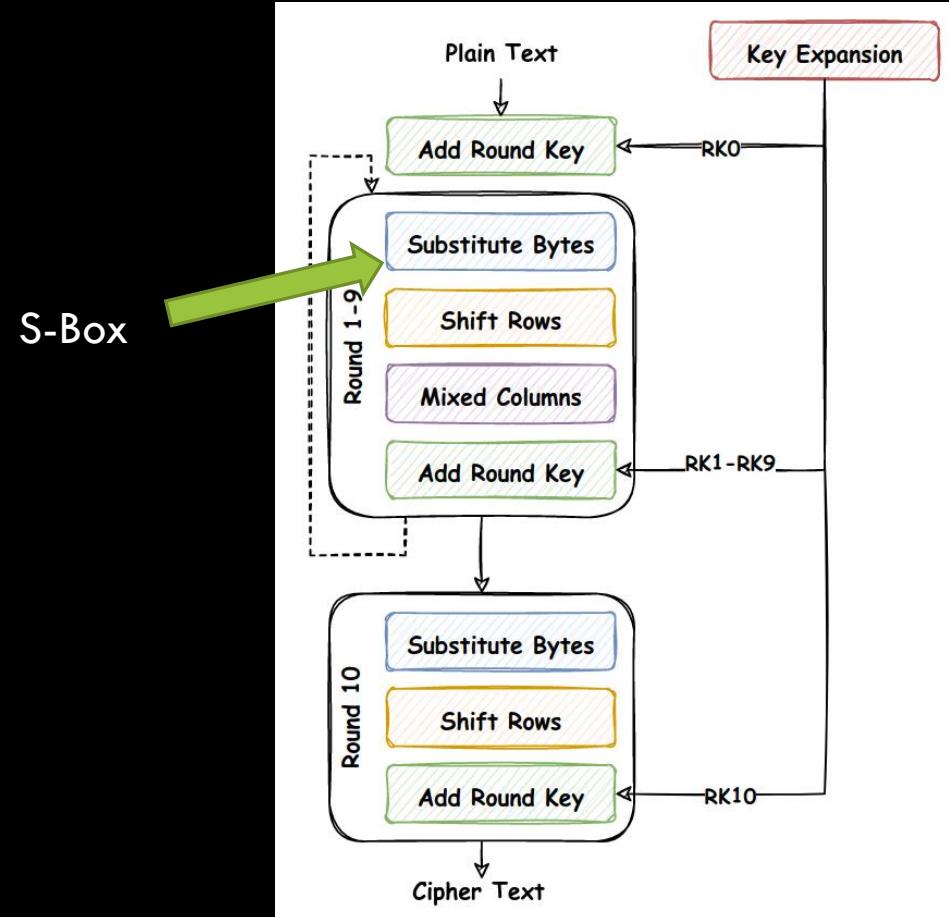
never need to be accessed by software. This is necessary to enable device bound encryption as the device unique key can only be used on the device. In addition to hardware keys, software key set can be stored in the internal memory of

AES in brief

- Advanced Encryption Standard (AES)
 - FIPS PUB 197: Advanced Encryption Standard (AES)
 - ISO/IEC 18033-3: Block ciphers
- Features:
 - Symmetric cipher
 - Block cipher: 128 bits (regardless of key size)
 - Keys: 128, 192 and 256 bits
 - Number of rounds depends on key size (10, 12 or 14 rounds)

Algorithm (128-bit key, encryption)

- Key Expansion
 - avoid using the same key each round
- Add Round Key
 - state \oplus RKn
- Substitute Bytes
 - apply S-Box to each byte of the state
- Shift Rows
 - Row bytes rotation
- Mix Columns



A “gentle” intro to... Side channel analysis (**SCA**)

*“A side channel is some observable aspect of a system
that reveals secrets within that system.”*

– The Hardware Hacking Handbook

Inception: Bell 131-B2 (1943)

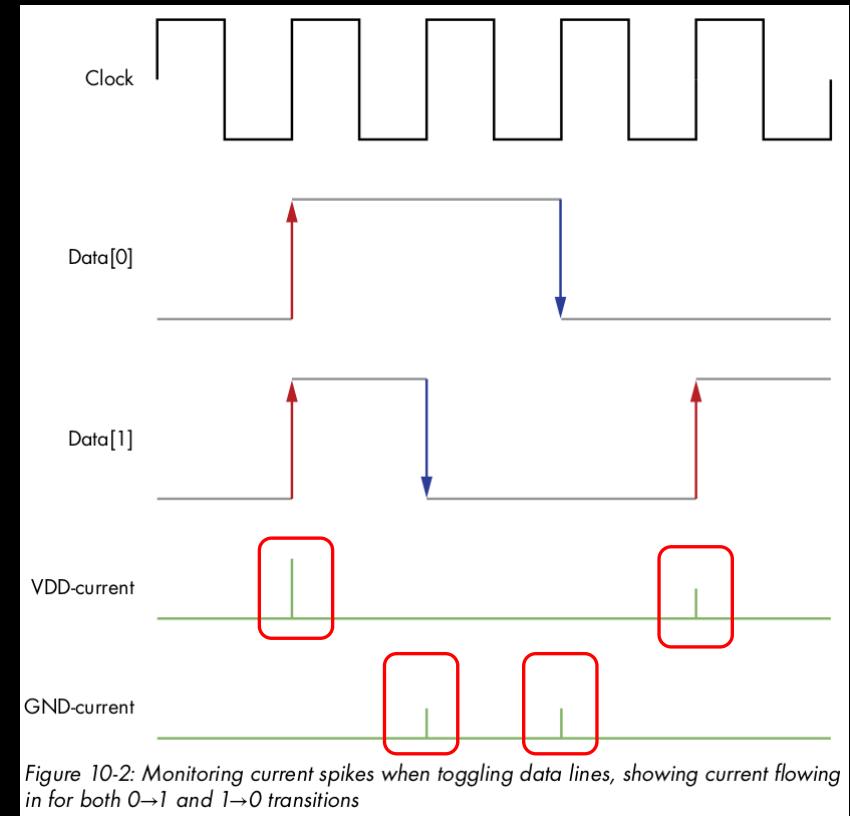
- Encrypted teletype used by U.S. Army and Navy:
 - One-time pads encryption
- Bell researcher noticed **spikes** in an **oscilloscope** nearby
- Disbelief: Is it really dangerous? **Prove** it!
 - Recovered 75% of plaintext from a different building (~25m away)
 - U.S. Army started clearing 30m perimeter
- Rediscovered in 1951
 - Recovered plaintext over power lines (400m away)



TEMPEST: A signal problem
NSA Declassified document

Power consumption

- Toggling data lines cause **current spikes** through VDD (+) and GND (-)
- **Energy** is a function of current flowing through circuit:
 - And so is Power!



Hardware Hacking Handbook
by Woudenberg & O'Flynn

Power leaks!

- Power consumption can leak information
 - E.g. during usage of (supposedly) secret data
- This includes intermediates of a cryptographic algorithm!

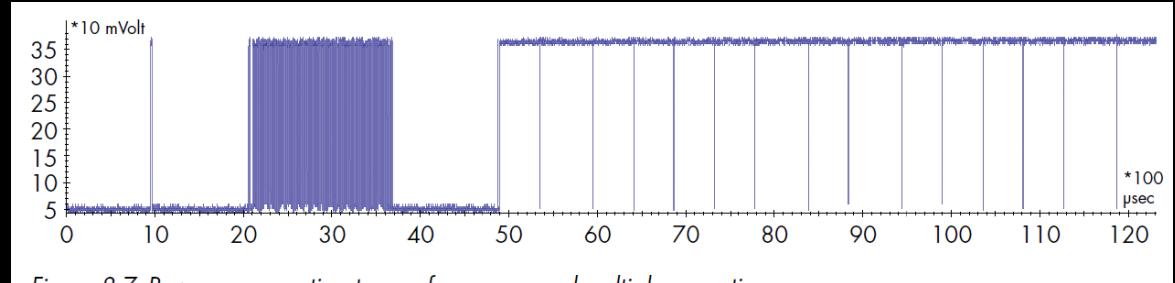


Figure 8-7: Power consumption trace of a square-and-multiply execution

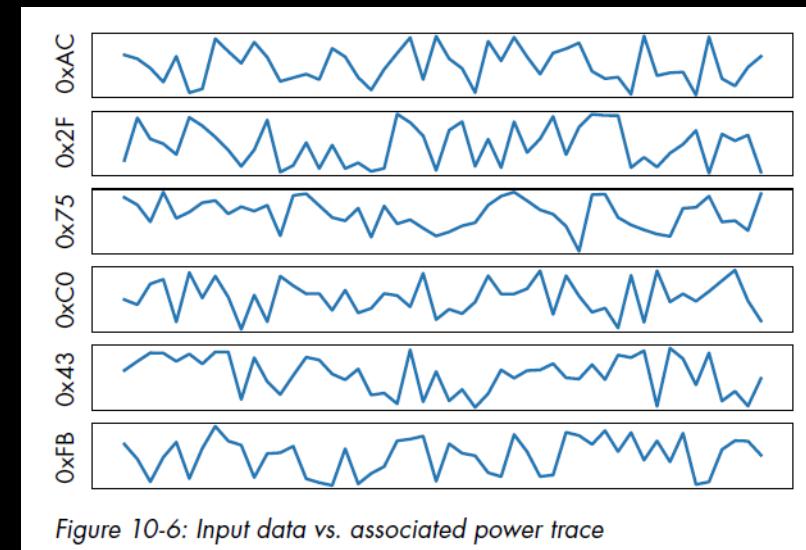


Figure 10-6: Input data vs. associated power trace

Hardware Hacking Handbook
by Woudenberg & O'Flynn

Electromagnetic field leaks too.

- Electromagnetic **emissions** can also leak information:
 - Program flow
 - Usage of (secret) data

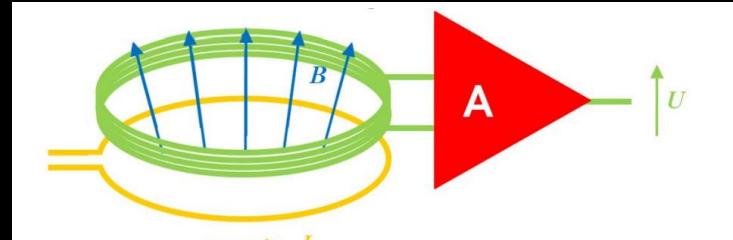


Fig. 4 Orientation of current loop (yellow) and sensor coil (green)

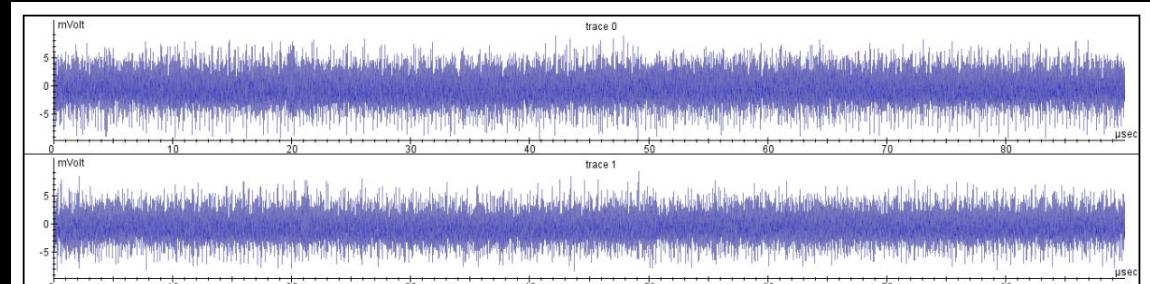


Fig. 12 Two unprocessed EM signal traces

Practical Electro-Magnetic Analysis
by Beer, Witteman, Gedrojc and Sheng

The challenge

- We have a device (**target**) performing **AES** encryptions
 - Using a HW cryptographic accelerator
- The key is **hidden** in HW
 - SW cannot access it
 - i.e. ANY code execution will not give you the key
- We can **encrypt**:
 - Whatever we want
 - As many times as we want

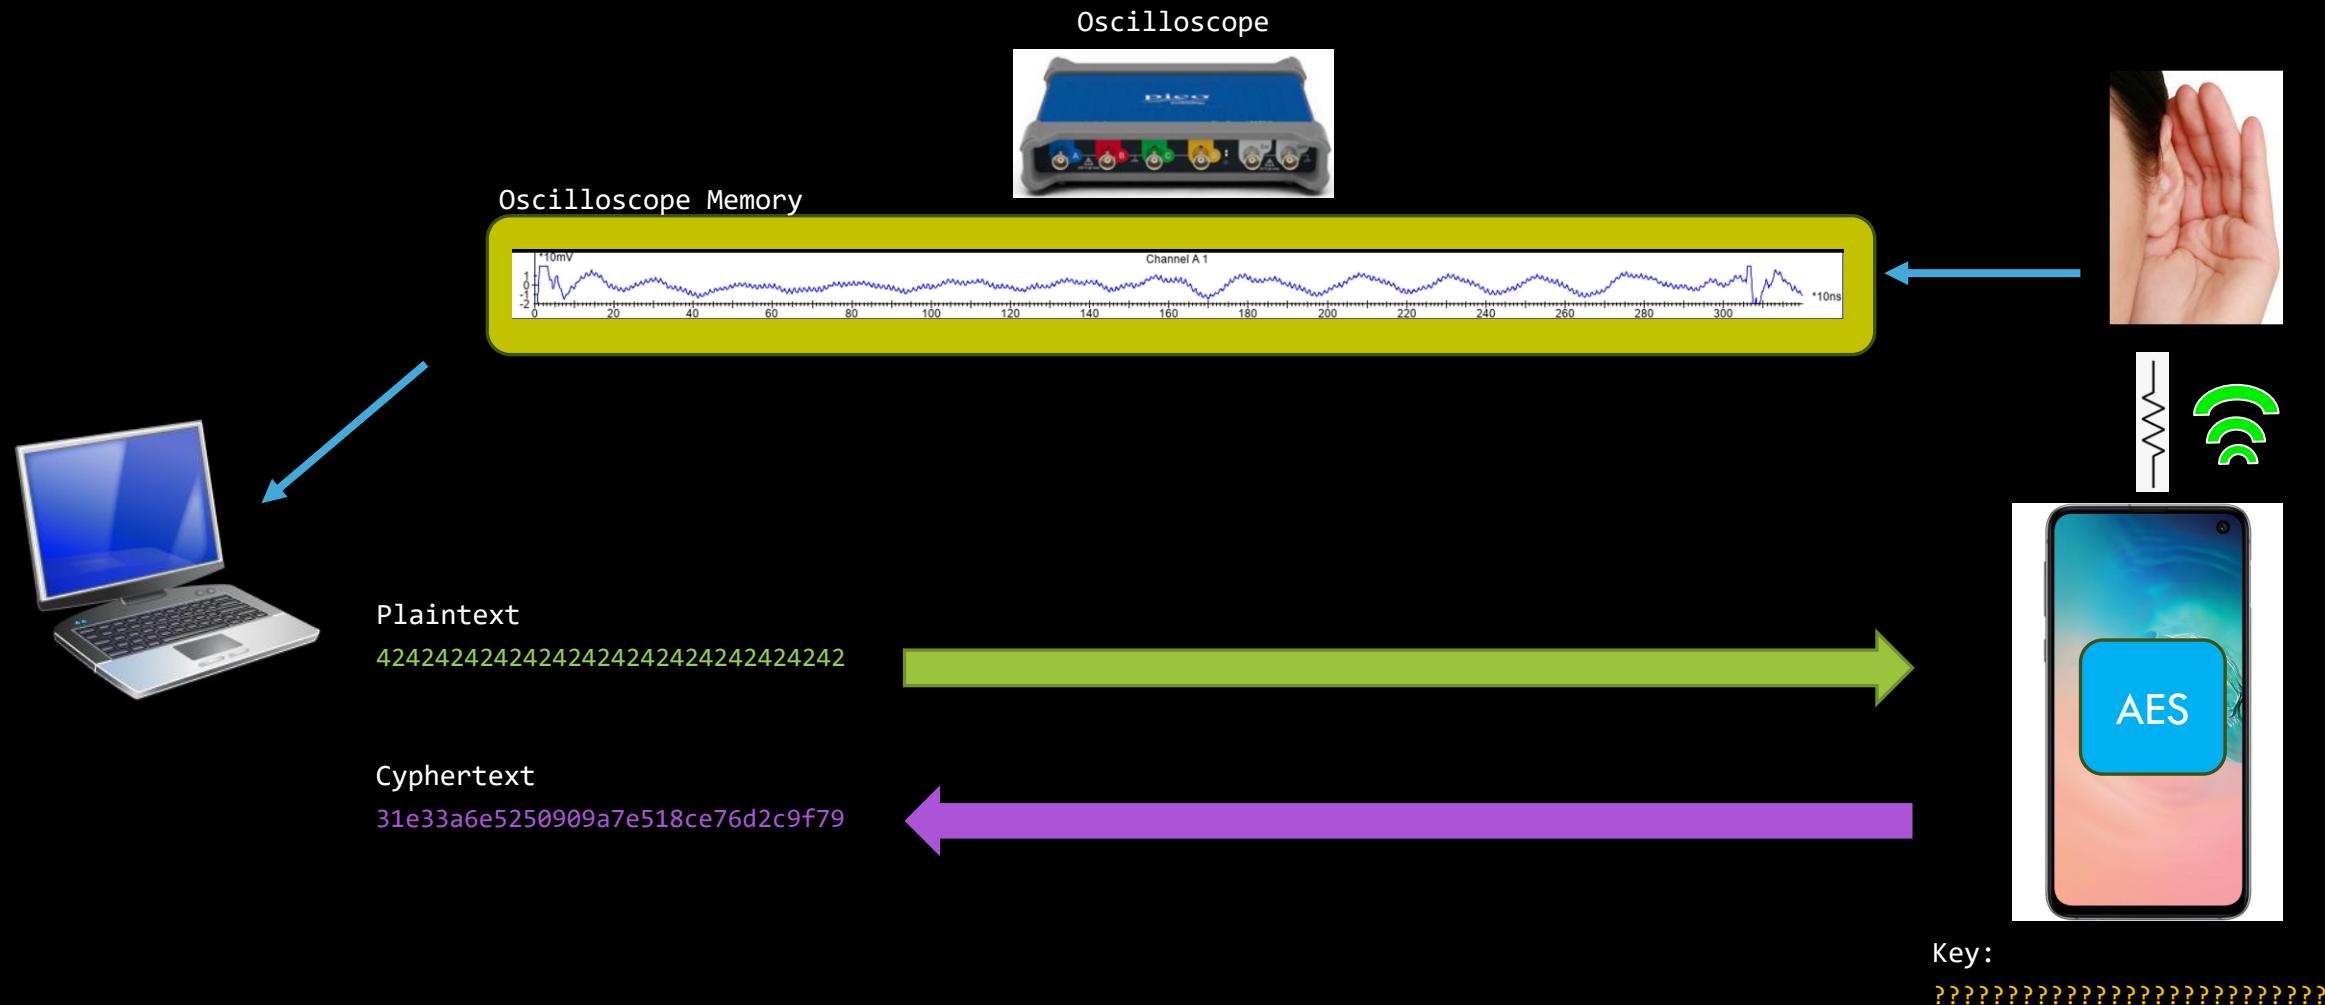
Can we **recover** the key?

Notes

- Next slides provide a **simplified** overview of Differential Power Analysis (**DPA**):
 - Well...we only have 45m for this talk 😊
- DPA is a renowned SCA technique:
 - Used in many labs around the world for security evaluations
 - Supported by many academic papers and...
 - ...many many **keys** extracted from real devices!

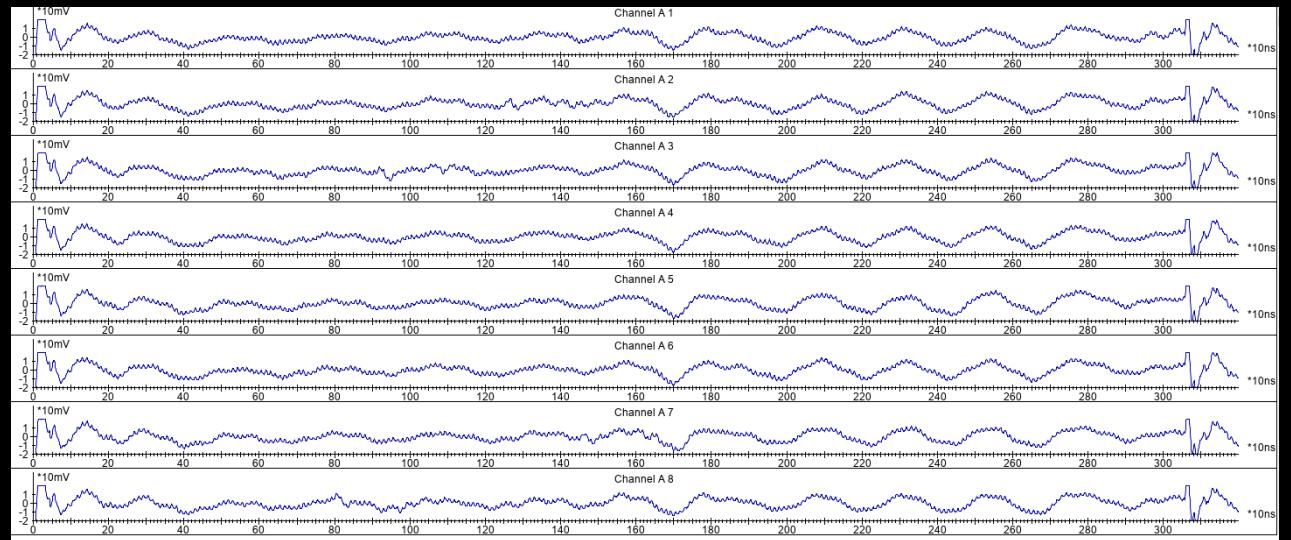
Feel free to ask for more info

The idea: Measure during encryption



The idea: Acquire MANY traces

- Execute the cryptographic algorithm a large number of times:
 - Vary the input **randomly**
- Acquire power traces:
 - while the cryptographic algorithm is being executed
- For each trace store:
 - The power **profile**
 - The input and the output (If available) **data**



Differential Power Analysis (DPA)

- Algorithm computes **intermediate values**:
 - Their actual value depends on the input and the key
 - Power profiles give information on the intermediate values
- Only **one** key can:
 - generate the right intermediates for all the input values
 - “**match**” the generated **power consumption profiles**

DPA: “Guessing the key”

- Select a key **candidate**
- Compute all **intermediates** for the candidate key:
 - For each input value
- Look for a key whose intermediates can “**match**” all power profiles for all input values
- In practice:
 - Compute **correlation** (Pearson) between the intermediates and traces values (at each sample)
 - Correct key should exhibit highest correlation

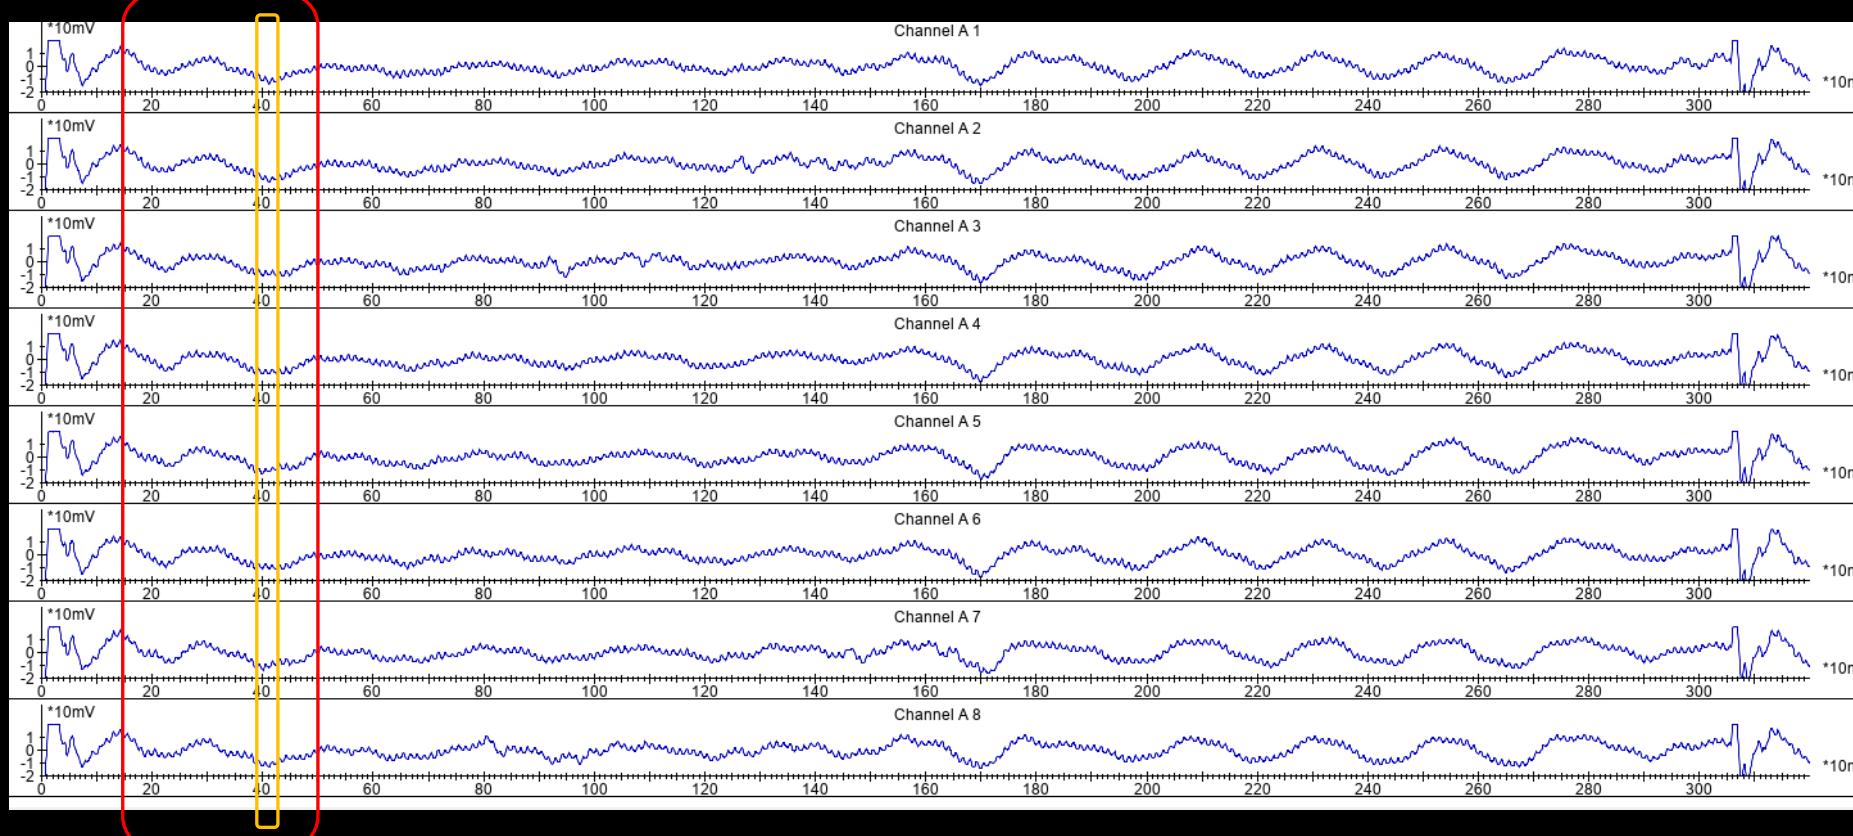
The idea: Matching

Key A matches!

Key candidate A
R1SubBytes[0][0]

Key candidate B
R1SubBytes[0][0]

7b



R1 SuBytes operation

6e

74

17

a1

7b

56

68

8a

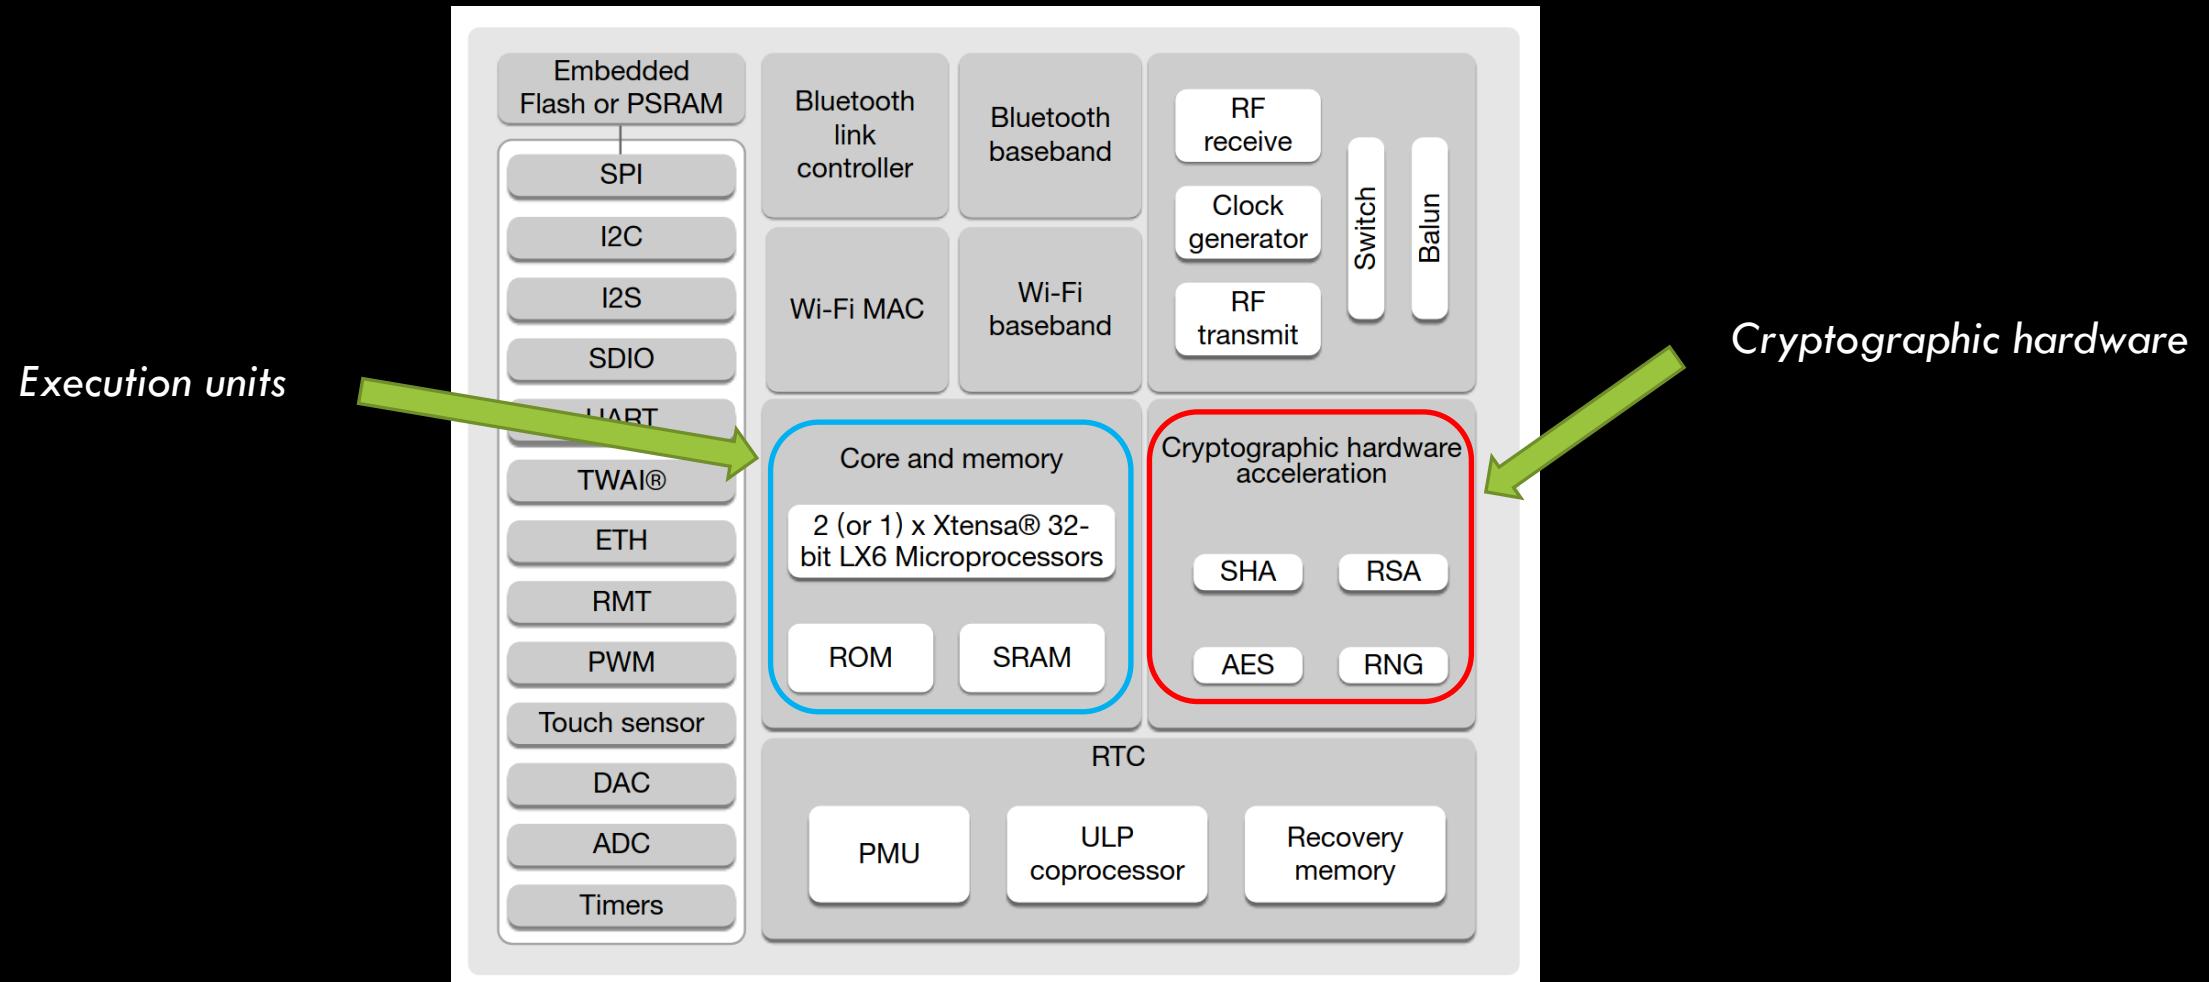
Our target: ESP32.

Espressif ESP32



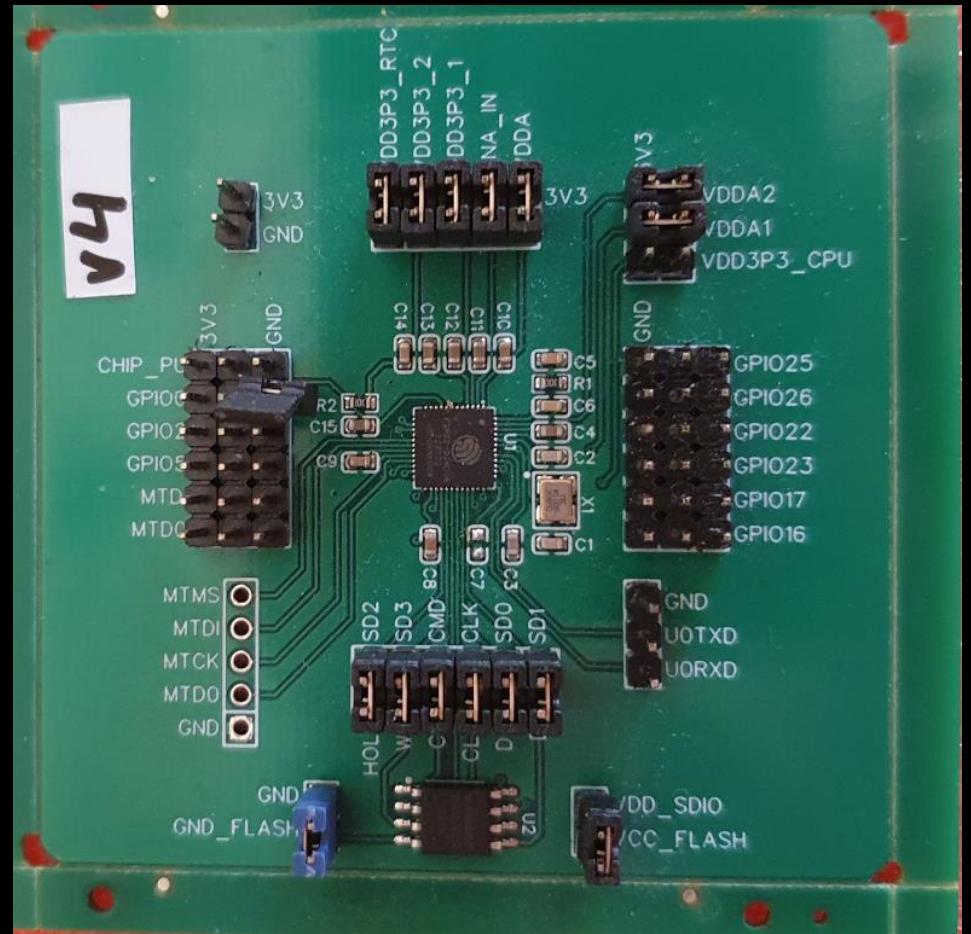
ESP32-D0WDQ6

ESP32 SoC: diagram



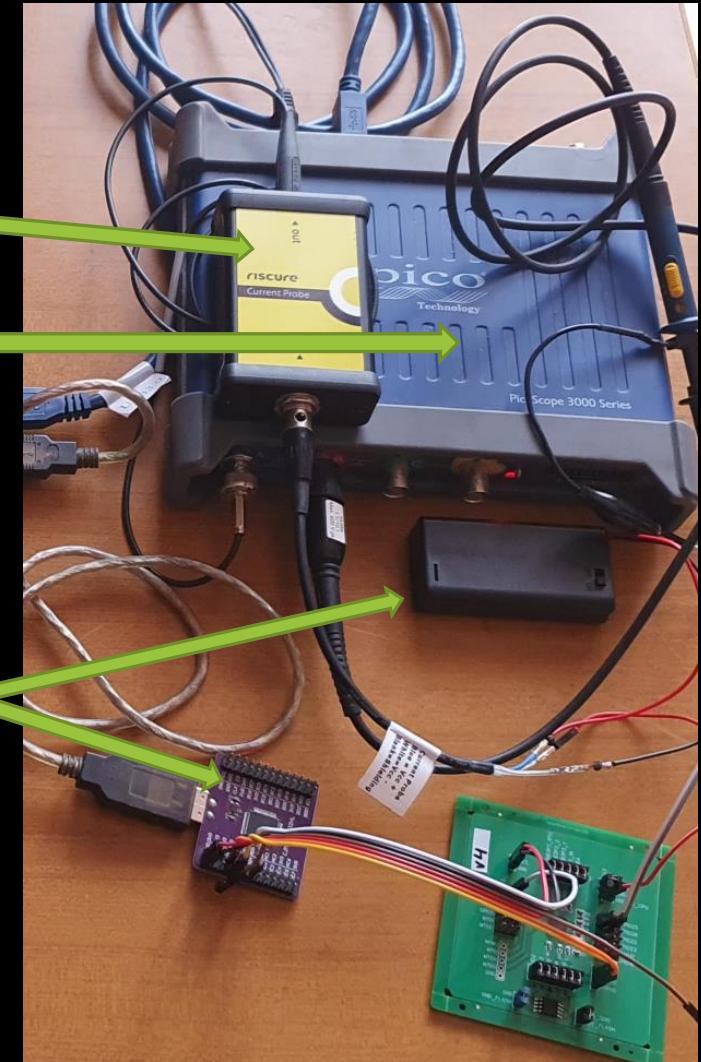
Our setup: Target

- Custom board:
 - Easier access to signals
 - Power CPU subsystem **independently**
- Application on target that can:
 - Set an arbitrary key
 - Operate the HW engine to perform **encryptions/decryptions**
 - Send a **trigger** to oscilloscope to start acquisition before encryption starts



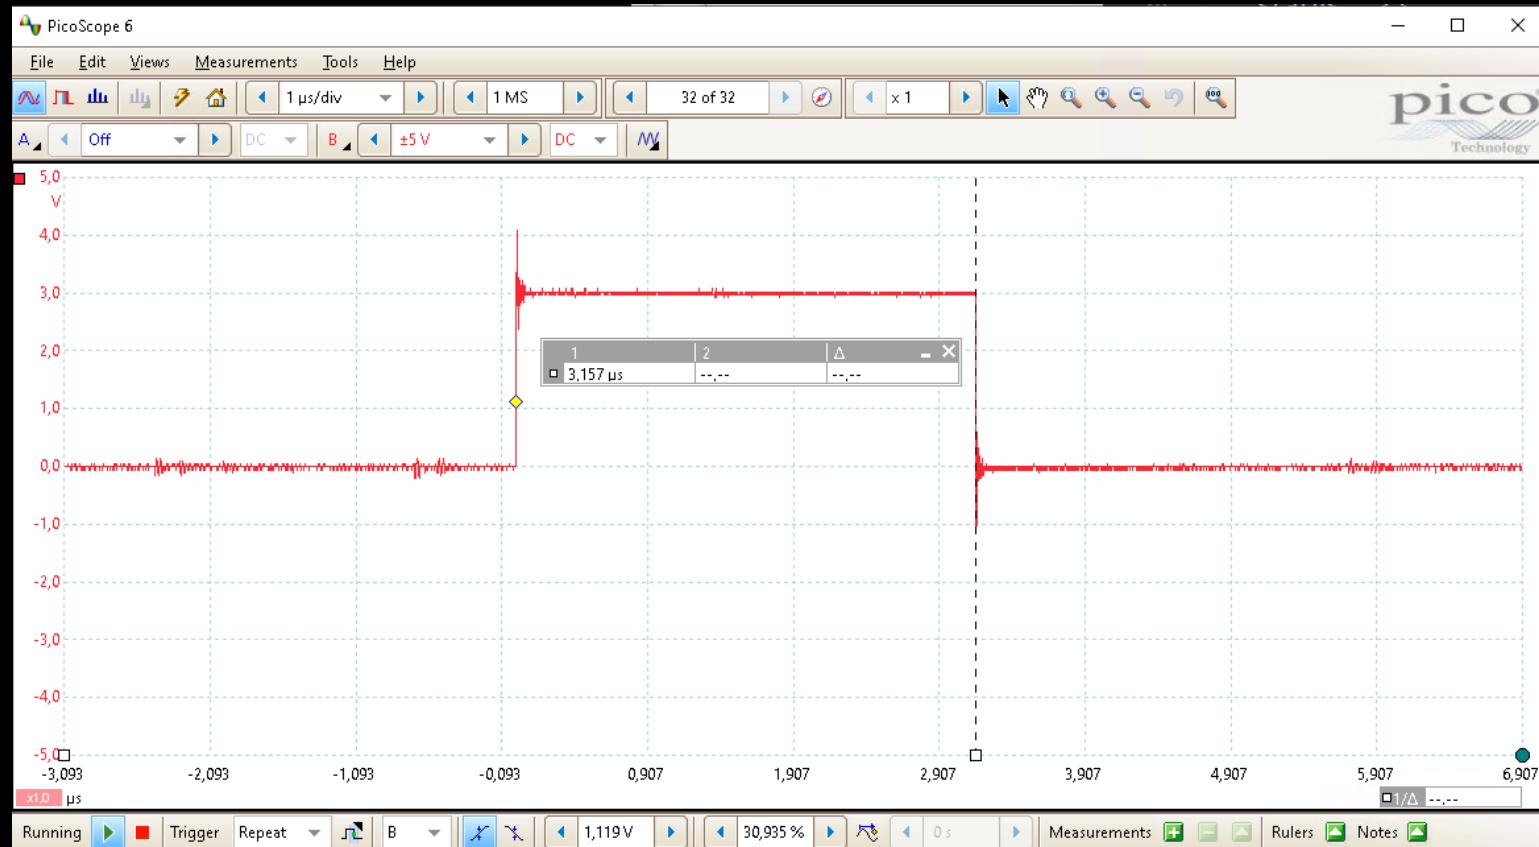
Our setup: Acquisition

- Riscure Current Probe:
 - For **power** measurements
- Picoscope 3406D:
 - Love that **scope**!
- FTDI 2232H:
 - **Serial** communications
 - For sending plaintext and receiving cypher text
 - **Power** the target: 3.3V
- A separate 3.3V battery package:
 - Cleaner measurements



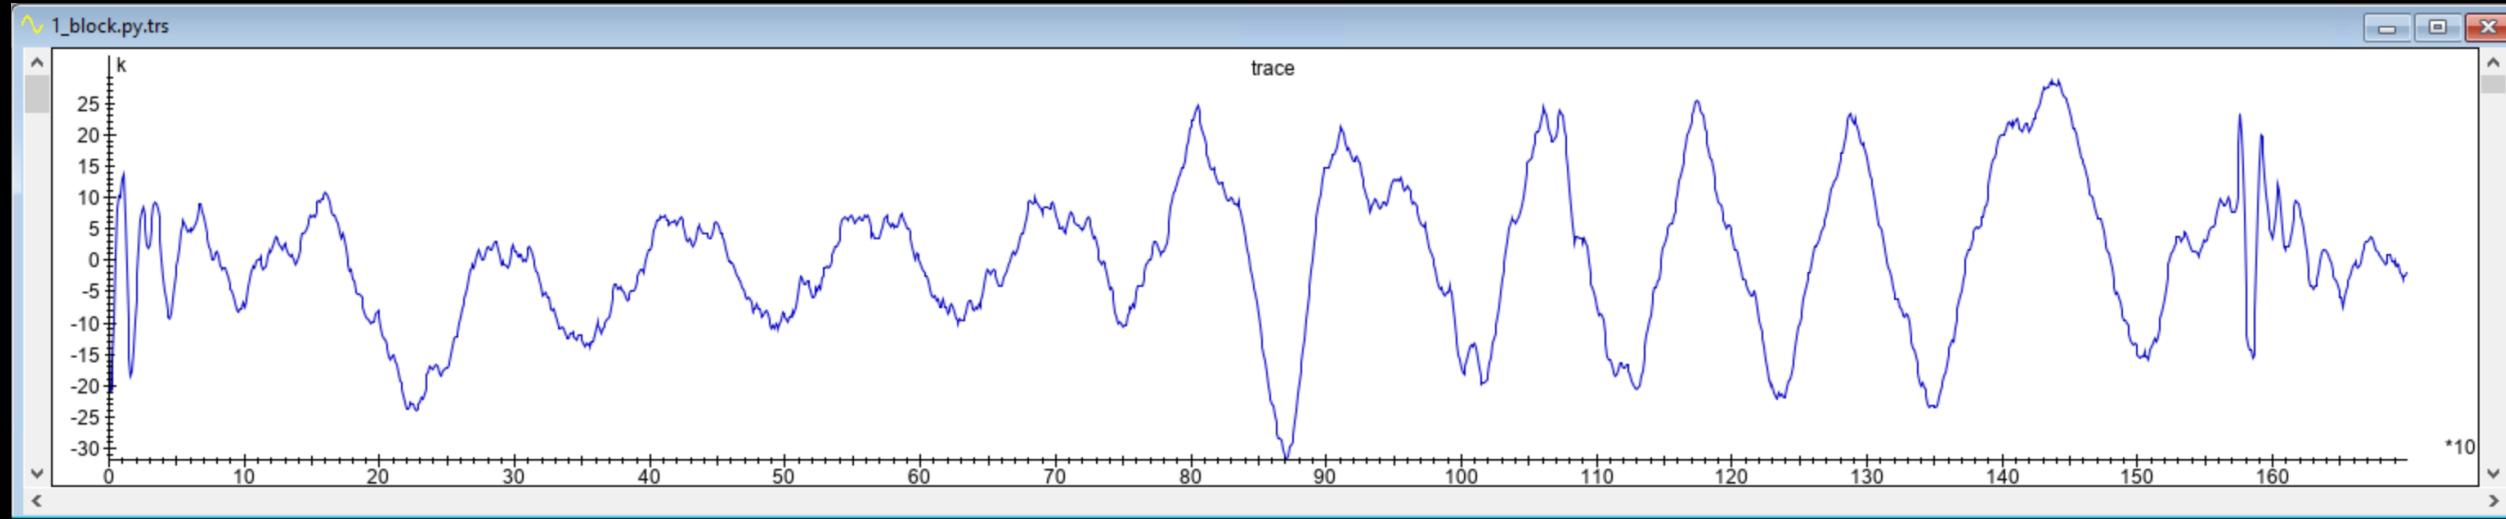
Reconnaissance and acquisition.

AES128 Encryption: Duration



- Time for one single encryption: 3.157 us

AES128 Encryption: Power profile



- Power profile does not show evident AES128 artifacts:
 - E.g. 10 repeated patterns (rounds)

Let's collect more traces!

- 20000 traces
- 1700 samples at 500 Mbit/s:
 - i.e. We acquires 3.4us
- Acquisition time: 3m 02s
- Acquisition speed: ~9.47 Million traces/day

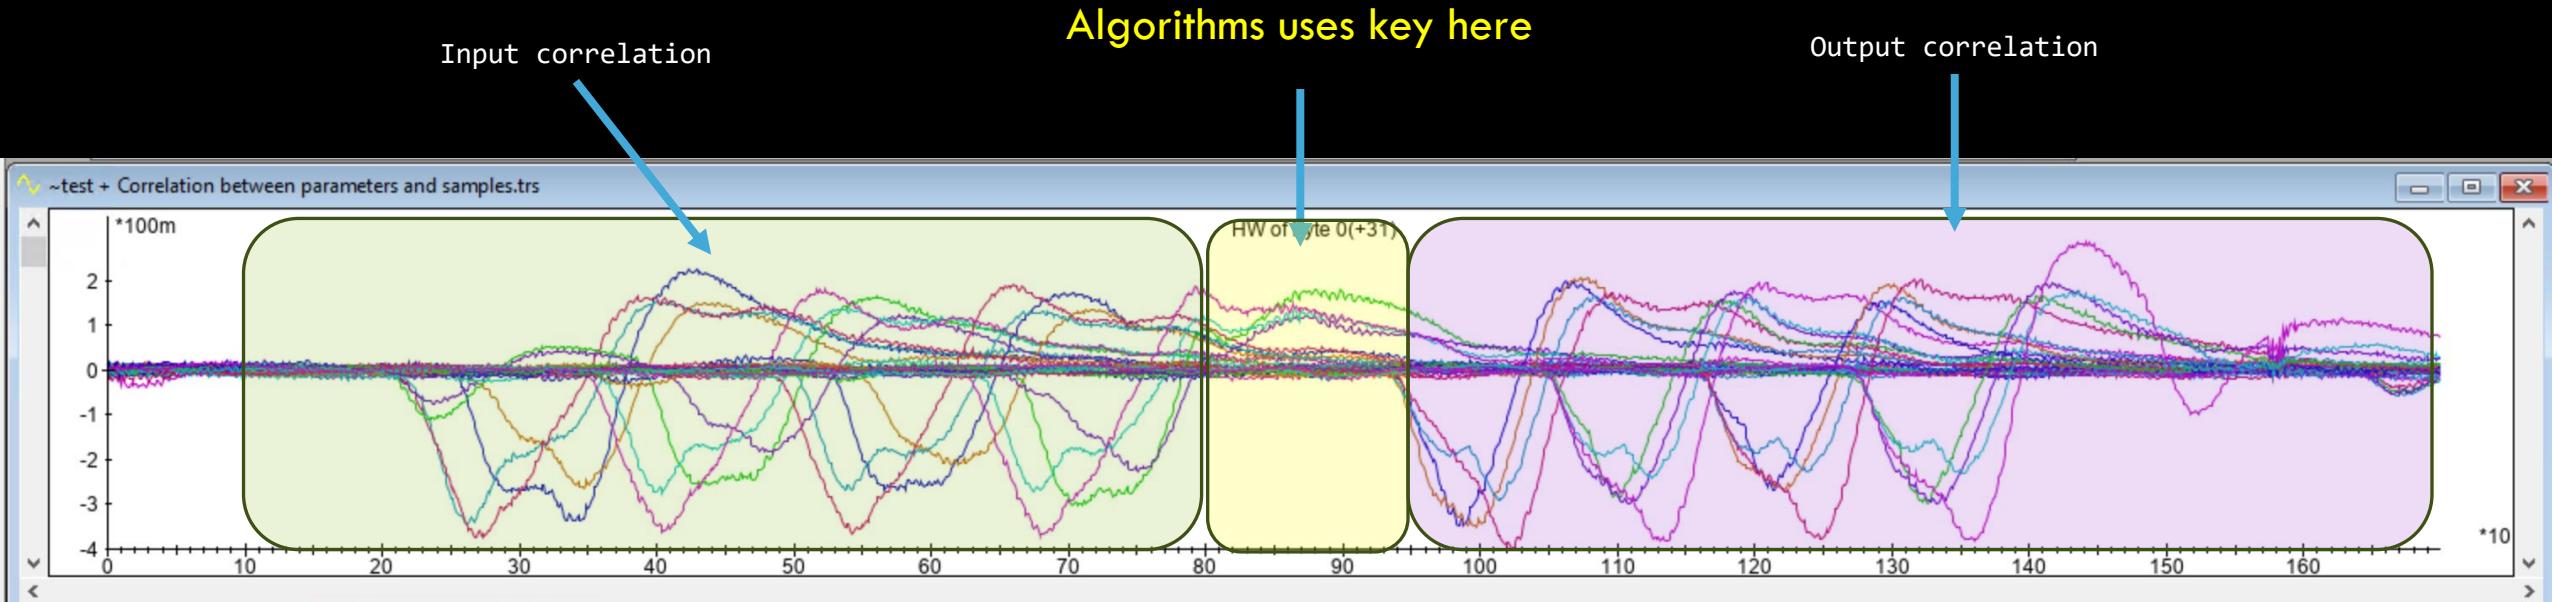
```
59 iterations = 1
60 nr_of_traces = 20000
61 delay = 0
62 nr_of_samples = 1700
Speed per day = 9.47E+06 (time: 0:03:02; traces: 19995)
Speed per day = 9.47E+06 (time: 0:03:02; traces: 19996)
Speed per day = 9.47E+06 (time: 0:03:02; traces: 19997)
Speed per day = 9.47E+06 (time: 0:03:02; traces: 19998)
Speed per day = 9.47E+06 (time: 0:03:02; traces: 19999)
Speed per day = 9.47E+06 (time: 0:03:02; traces: 20000)
20000
[Finished in 185.3s]
```

Can we get the key?

Considerations: Speed

- HW cryptographic engines can be **fast**:
 - Lower number of **samples** required (w.r.t SW implementations)
 - Operation completes in a shorter **time** (w.r.t SW implementations)
- Acquiring **Millions of traces/day** is not uncommon
 - Even with **very simple setups**

Input/Output correlation



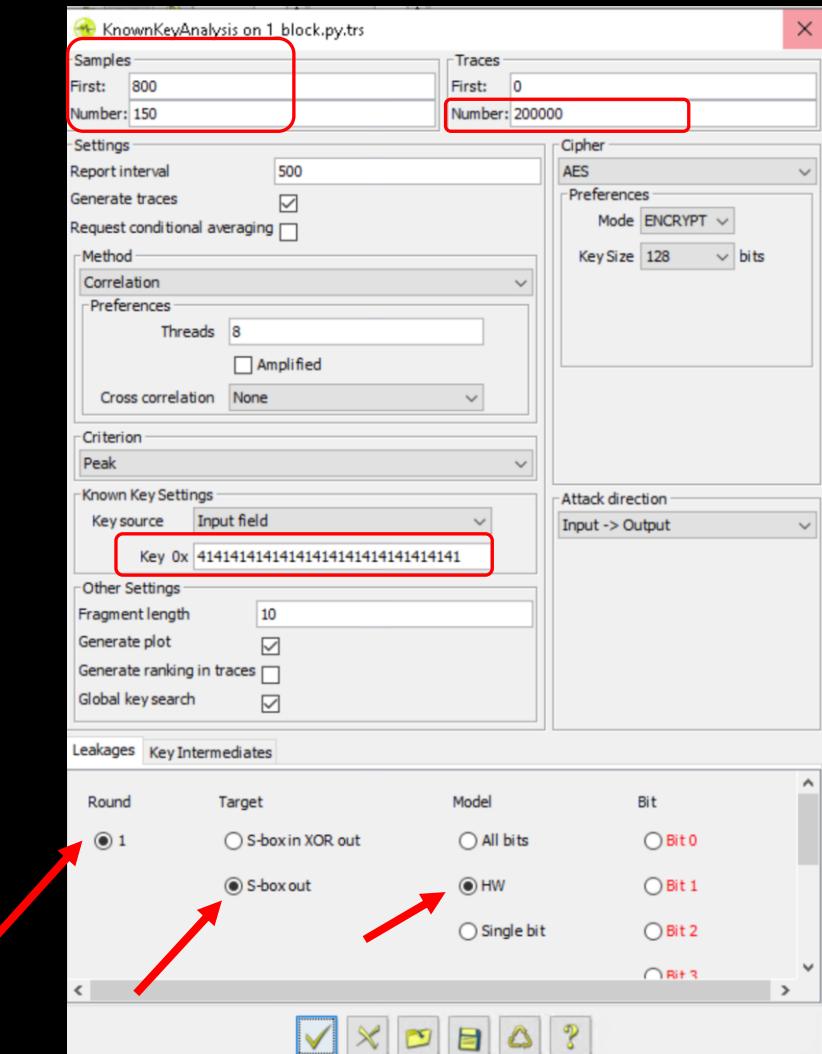
- Correlation with input (plaintext) and output (ciphertext) bytes:
 - Shows where such bytes are being “used”
- Key used:
 - after input is received
 - before output is generated
 - between samples 800 and 950

Known key analysis

- Take **another** device identical/similar to your target:
 - Same SoC
 - Configuration as close as possible
 - You must control it (i.e. be able to set your key)
- Set your **own** key → You can compute intermediates (for every input)
- Perform correlation analysis with power profiles
- You will get:
 - If the SoC **leaks** information
 - **Where** information leakage happens

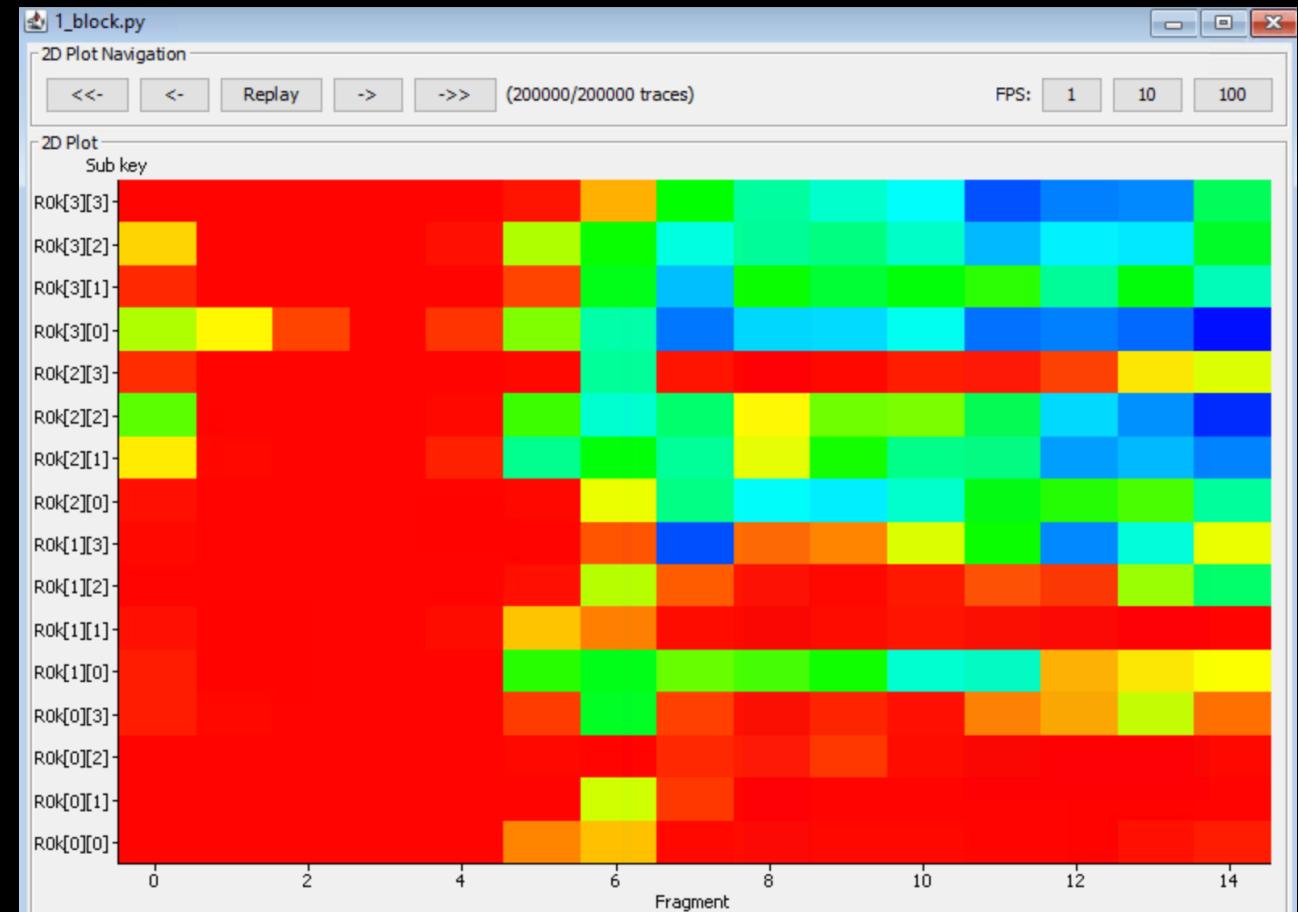
Known key analysis: Settings

- Performed on 200k traces:
 - Acquired in 30-40m
- Focus only between samples 800 and 950
- Leakage model:
 - Hamming weight on S-box output



Known key analysis: Results

- Leakage for all key bytes:
 - Samples: 820 → 840

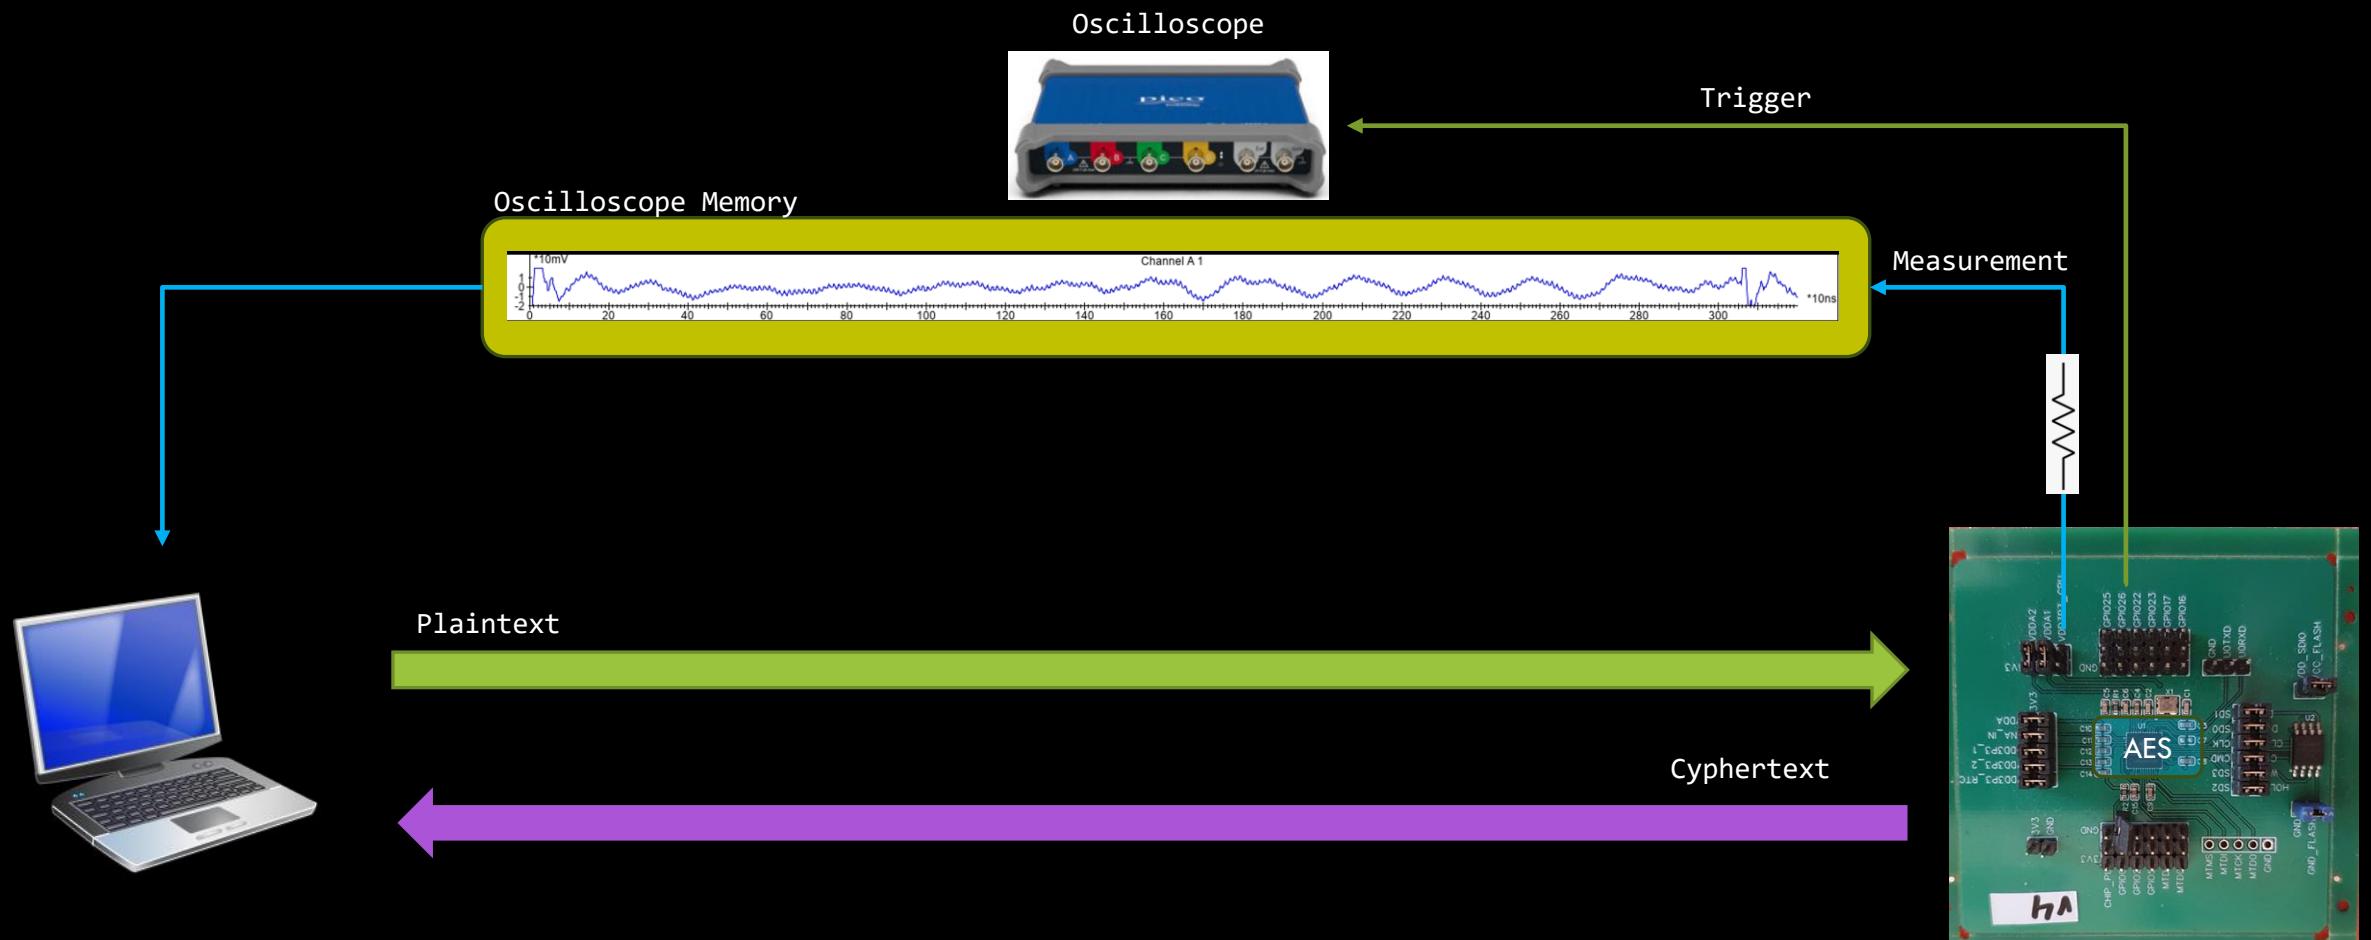


Success!

- Key can be retrieved in ~40m:
 - 200k traces
 - Acquisition time: ~30m
 - Acquisition speed: ~9.4M traces/day

Can we go faster?
Segmented memory.

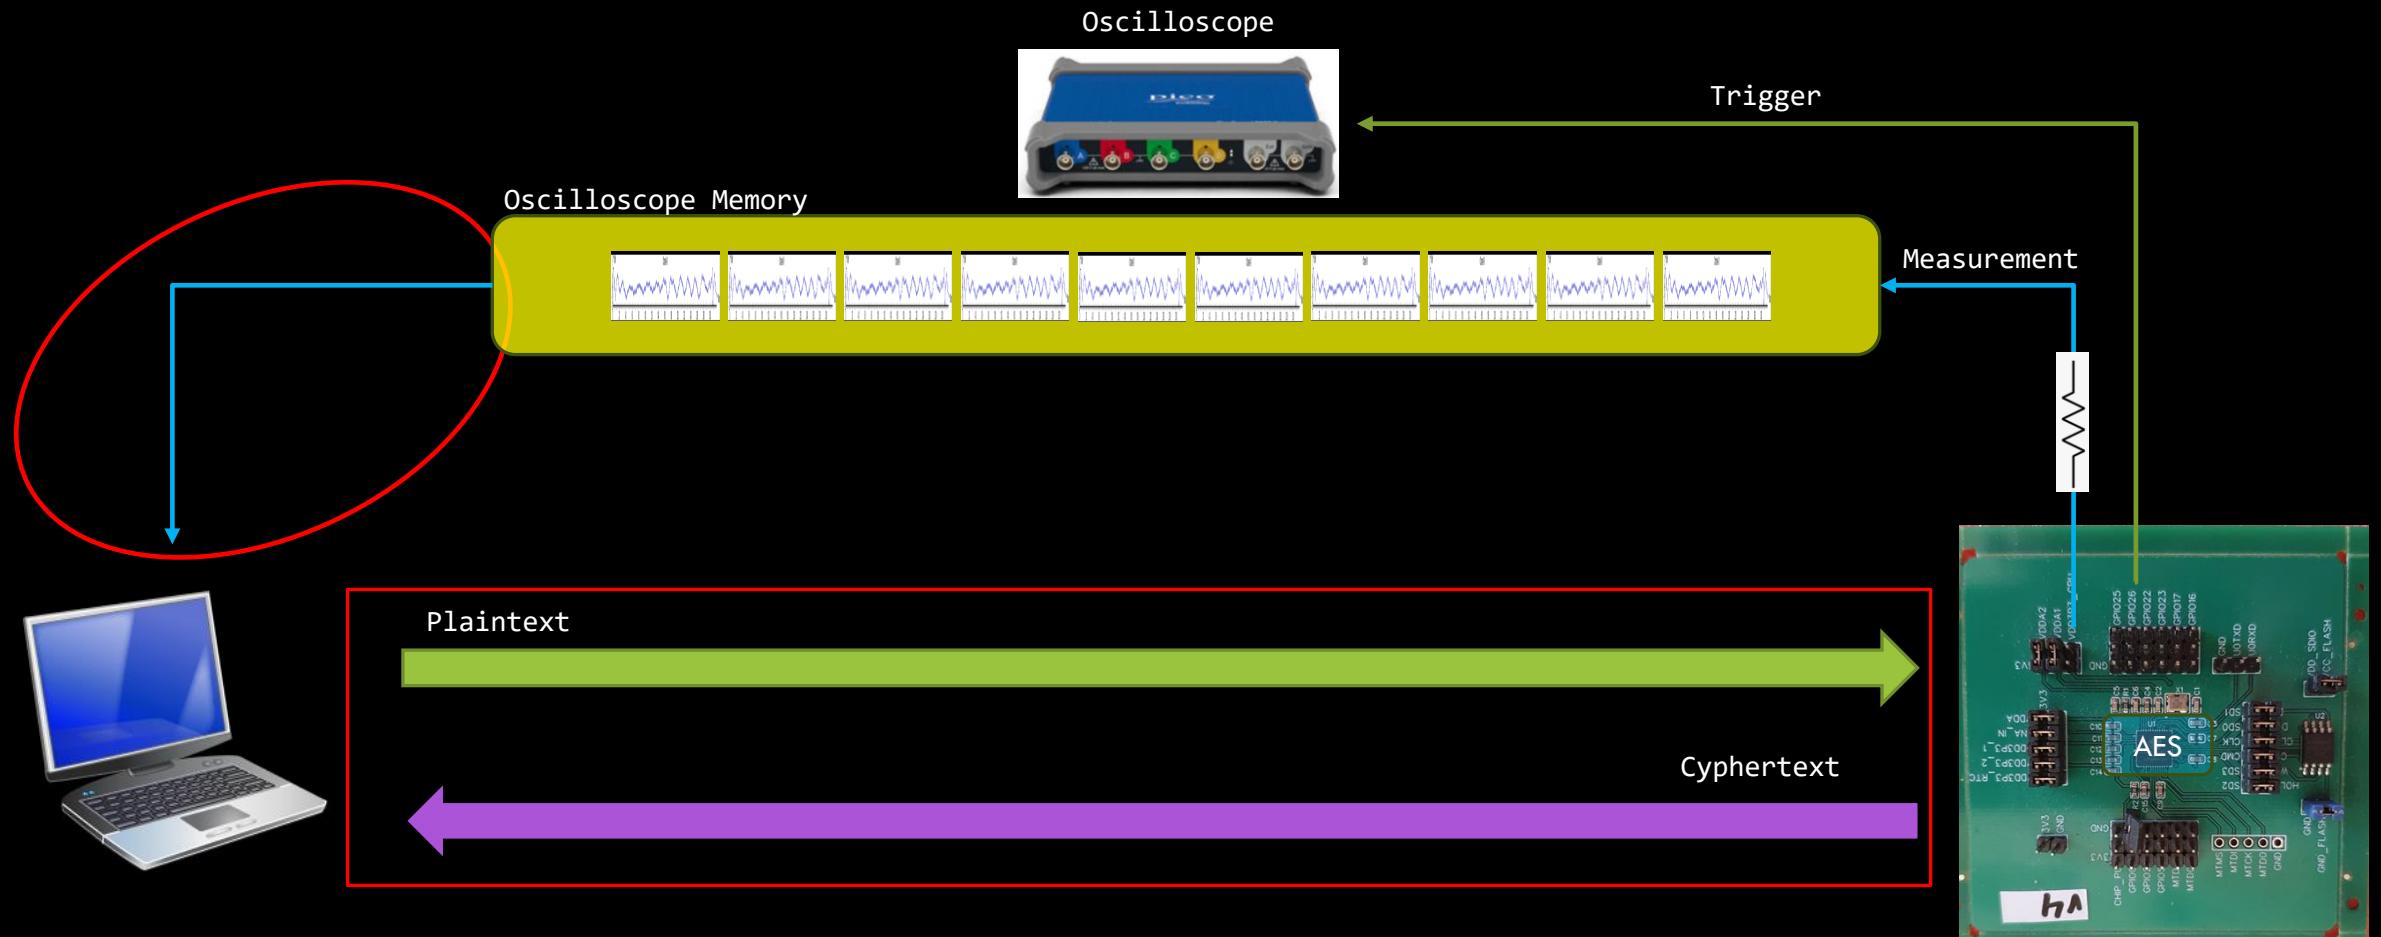
Acquisition cycle



Segmented Memory

- Feature available on many modern oscilloscopes
- Scope internal memory can be “segmented” to store **multiple** traces
 - Number limited by scope memory size
- Acquired traces are sent to PC in one single **bundle**
- Typical usage:
 - Perform multiple measurements with the same **input**
 - Traces can be averaged to reduce noise

Acquisition: Segmented memory



Segmented memory + averaged traces

- 20000 traces:
 - 100 traces in segmented memory (iterations)
 - 2000 input provided
- 150 samples
- Acquisition time: 25s
- Acquisition speed: ~675 Million traces/day

```
29
30 iterations = 100
31 nr_of_traces = 2000
32 delay = 800
33 nr_of_samples = 150
34
Speed per day = 6.75E+08 (time: 0:00:25; traces: 199200)
Speed per day = 6.75E+08 (time: 0:00:25; traces: 199300)
Speed per day = 6.75E+08 (time: 0:00:25; traces: 199400)
Speed per day = 6.75E+08 (time: 0:00:25; traces: 199500)
Speed per day = 6.75E+08 (time: 0:00:25; traces: 199600)
Speed per day = 6.75E+08 (time: 0:00:25; traces: 199700)
Speed per day = 6.75E+08 (time: 0:00:25; traces: 199800)
Speed per day = 6.75E+08 (time: 0:00:25; traces: 199900)
Speed per day = 6.75E+08 (time: 0:00:25; traces: 200000)
Traces in traceset: 2000
[Finished in 28.5s]
```

Can we get the key?

Nope.

- Input not sufficiently diversified
 - Only 2000 plaintexts
- No sufficient leakage to reveal key:
 - On this specific target
- How can we:
 - Have sufficiently **diversified** input AND
 - Minimize communication **overhead** with target

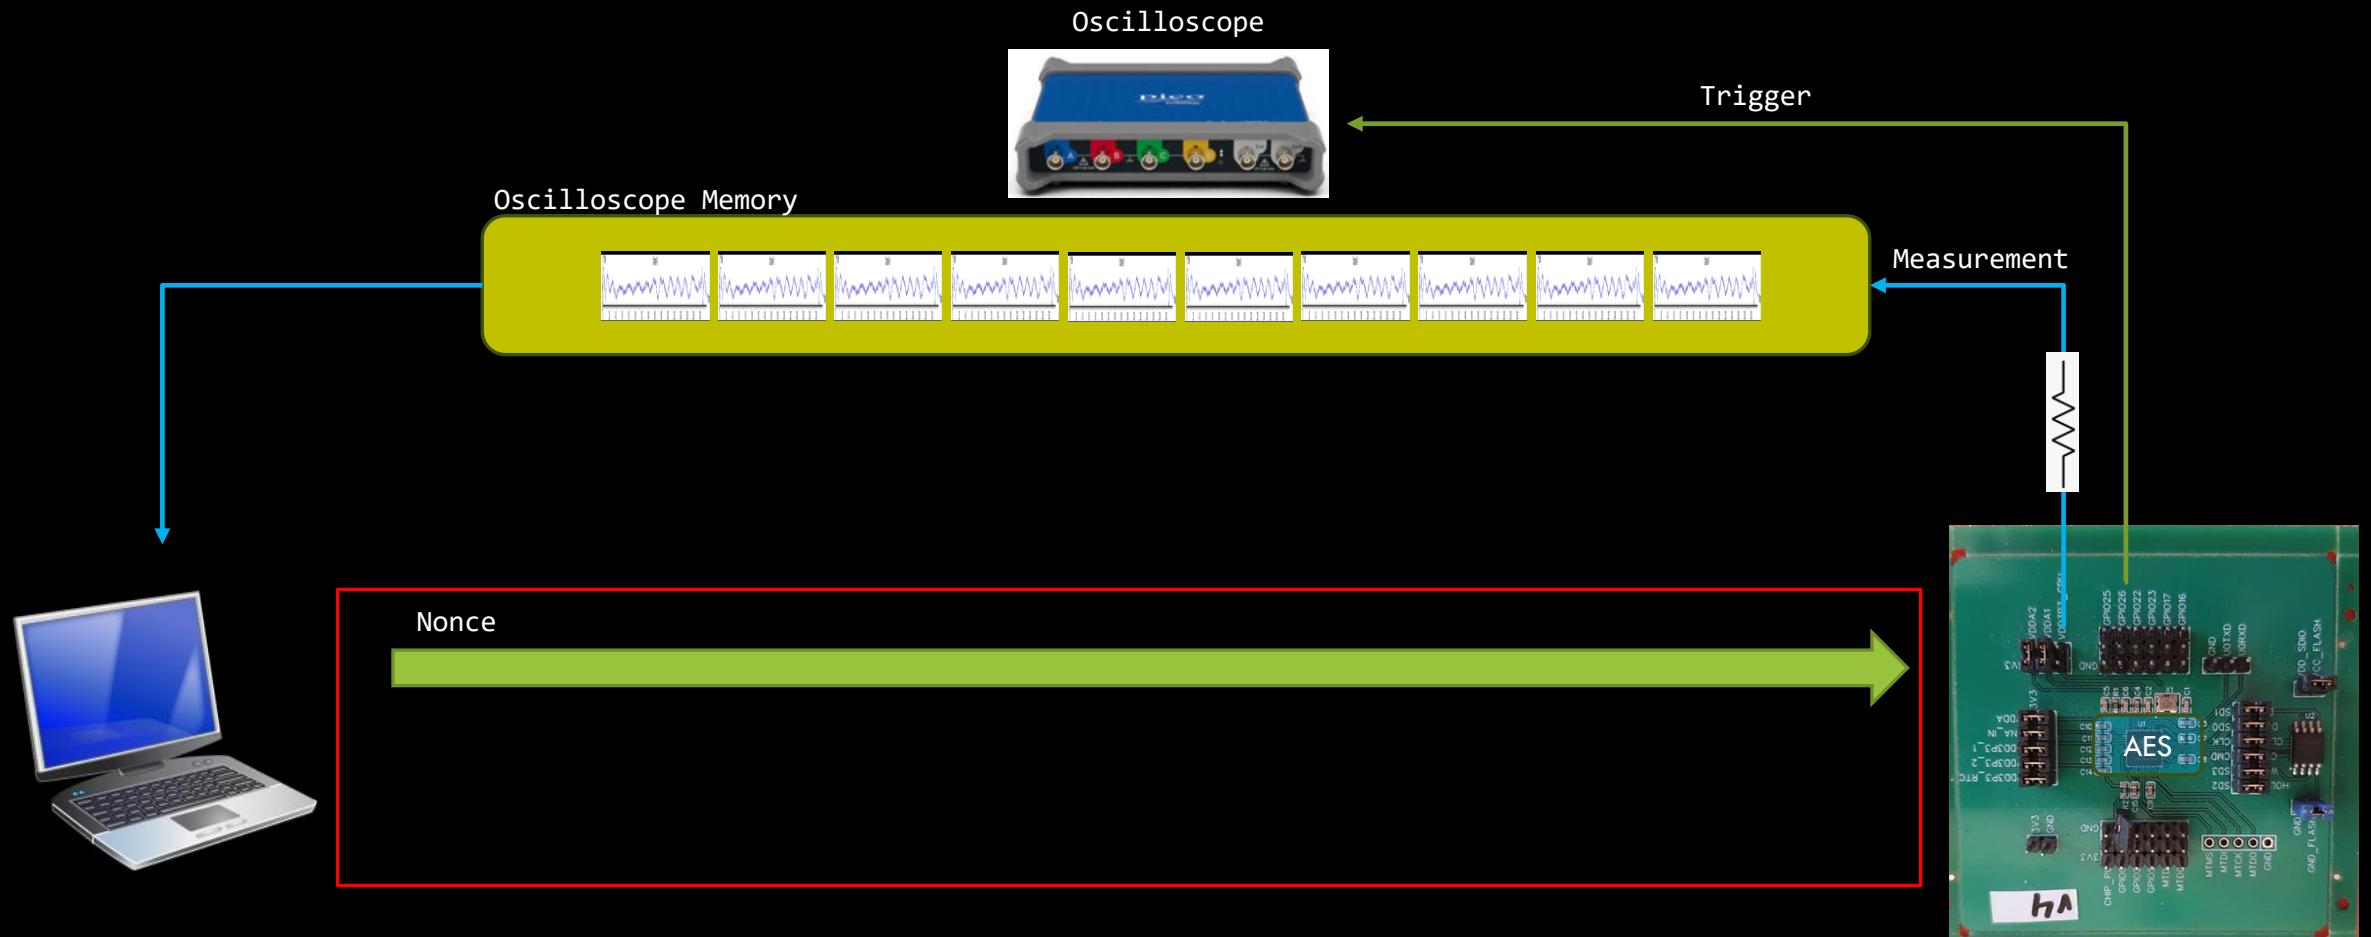
?

Speedy Gonzalez: On-target input generation.

Generating input on target

- No need to provide input from PC
 - It's only sufficient to **KNOW** the input to the AES engine for each encryption operation
 - Output not needed
 - We are attacking encryption → AES round 1
- General idea:
 - Only send an initial nonce
 - Use a cryptographic function to generate next input
 - Apply the same function on the PC side to compute the same input
- Examples:
 - Recursively apply a Hash function to nonce
 - Use AES engine output as input for next operation

Acquisition: On-target input generation



How fast can we go?

- 20000 traces:
 - 50000 traces in segmented memory (iterations)
 - 4 bulk transfers to PC
- 1 sample: (nr. 820)
- Acquisition time: 12s
- Acquisition speed: ~2 Billion traces/day

```
34 iterations = 50000
35 nr_of_traces = 4
36 delay = 820
37 nr_of_samples = 1
38
Speed per day = 1.78E+09 (time: 0:00:02; traces: 50000)
Speed per day = 1.92E+09 (time: 0:00:04; traces: 100000)
Speed per day = 1.97E+09 (time: 0:00:06; traces: 150000)
Speed per day = 2.00E+09 (time: 0:00:08; traces: 200000)
Traces in traceset: 200000
[Finished in 11.9s]
```

Can we get the key now?

Yes!

- Key can be retrieved in ~1.5:
 - 200k traces
 - Acquisition time: ~12s
 - Acquisition speed: ~2B traces/day

Notes

- Even faster acquisition speed may be possible with further **tuning**
- We can now retrieve a key from ESP32:
 - Used by the HW crypto engine
 - By means of power analysis
 - Using segmented memory
 - Generating input on target
 - Using JlscA for analysis
- In less than 25s.

Demo.

Technique is known and used!

- Research:
 - Leakage Assessment Methodology - a clear roadmap for side-channel evaluations - Schneider et. Al
 - A flexible leakage trace collection setup for arbitrary cryptographic IP cores - Moschos et. al.
 - Apple vs. EMA: Electromagnetic Side Channel Attacks on Apple CoreCrypto - Haas et. al.
 - Using a magic wand to break the iPhone's last security barrier – tihmstar
- Also security/evaluation labs are (likely) using it ☺

Back to base:
Conclusions.

Summary

- Acquisition speed of millions of traces/day are **common**
- **Billions** of traces/day can be achieved:
 - Under specific conditions
 - Some degree of target control is required
- Technique is known, described in literature and actively used

Implications

- Claims of “resistance to SCA” should consider acquisition speeds in the billions of traces/day:
 - Evaluation/security labs
 - Certification schemes
- Very fast attacks may be possible in some specific scenarios:
 - E.g. when access to target is time constrained

raelize

Thank you!

Cristofaro Mune
cristofaro@raelize.com
[@pulsoid](https://twitter.com/@pulsoid)